

UMBRAL

Business Bending Packet / First-Proof / OPERATOR

Functional educational demonstration. Team Blueprint draft status preserved. Live LLM activation and original recording evidence remain outstanding.

Rodrigo Peña León - OPERATOR - 7 October 2026

1. Decision and evidence status

The supplied Team 2 Blueprint selects FIRST-PROOF as its primary vacuum. This individual build implements the OPERATOR declaration: intake-to-review workflow with scope, reviewer, deadlines and consent. The source remains a draft synthesis of four role briefs, with proposed bets and a pending TECHNOLOGIST declaration. Starting this build is Rodrigo's instruction; it does not establish a completed five-person vote, live Team session, employer commitment or real pilot result.

The personal Brain brief remains Rodrigo's final individual brief and preserves his preference for paid Micro-experience. This build follows the supplied Blueprint's First-Proof direction while retaining that dissent: a verified simulation may fail to create income, and employer task supply/reviewer capacity may make the model uneconomic.

This Packet was saved before product-source code. Any later scope change must be dated here.

2. Exact user and failed moment

Primary user: a coordinator at a Mexican university career office or small-business network. Failed moment: a candidate with informal bookkeeping experience cannot show a credible, bounded sample, while the employer has no named reviewer, rubric, agreed review deadline or permission to use/share the evidence. These are hypotheses derived from the Blueprint, not interview findings from this build.

Secondary actors: the candidate who owns the evidence, and the named human reviewer who observes the explanation or live changed-condition exercise and records feedback. The coordinator owns delivery, not an employability verdict.

3. Product and selected slice

UmbraL is a controlled workflow for one fictional bookkeeping work sample. It demonstrates ability in the bounded task; it is not a universal certificate. The first viewport is an operator work queue with session states, review deadlines and a create-invitation action.

Included: invitation with bounded scope and fictional/public data; named reviewer and deadlines; candidate assistance disclosure; observed changed condition; criterion-by-criterion human review; candidate-first result; correction request and reviewer response; private approval; individual recipient authorization and revocation; local export; audit trail; review-time measurement; pilot decision thresholds.

Excluded: CV builder, job board, ranking, automatic rejection, universal score, real production work, candidate fees, credential verification, identity guarantee and claims of hiring. The demo has a role switch for exploring fictional records, not real multi-user authentication. Do not enter personal/client data. Local persistence is explicitly limited to this demonstration on this device. Production needs authenticated actor isolation, retention/deletion policy, a real consent service and institutional review before real participants.

4. One coherent workflow

1. Operator chooses the verified fictional template, names the reviewer, sets due dates and confirms task scope/review conditions.
2. Candidate sees the task, rubric, data and assistance policy before accepting. An accessible written explanation alternative is available.
3. Candidate submits an explanation and discloses AI assistance. The reviewer observes a live revision/explanation of the changed condition.
4. Reviewer records evidence per criterion, observation method, feedback and review time. No automatic rating/ranking. Incomplete observation blocks a shareable proof.
5. Result returns privately to the candidate. They can request correction, keep it private, remove their local attempt or approve the accurate record.
6. Candidate authorizes a named recipient for a stated purpose, then exports a recipient-specific proof. Authorization can be revoked locally; already downloaded copies cannot be recalled. Never imply otherwise.

State rules: invited > submitted > reviewed_private > correction_requested > reviewed_private > approved_private > authorized. A private incomplete attempt may be revised or deleted; it is not exported as a permanent failure label. A reviewer must resolve a correction before candidate approval/sharing. An edit to a reviewed record invalidates prior approval and sharing. A proof always carries fictional-demo labeling.

5. Dragon Stack

LLM: a server-side drafting assistant for routine wording of scope, accommodation and rubric, grounded in the selected template. It must not evaluate a candidate, rank, reject, verify identity or infer employability. Provider key remains server-side. Without a configured key, show that AI is unavailable and let the verified template workflow continue. Do not label the deterministic fallback as LLM output. Live LLM activation is a delivery dependency until a valid provider secret and live call are verified.

Structured/verified data: versioned template JSON, explicit criterion definitions, synthetic ledger rows, a manually checked reference calculation, provenance, allowed state transitions and consent/event records. These data establish the fictional exercise arithmetic, not actual skill or demand.

Additional component: deterministic workflow/consent guardrails, deadline/review timer and portable recipient-specific evidence export. The browser performs these functions independently of a model.

Implementation: semantic HTML/CSS and browser ES modules; a minimal Cloudflare-compatible Worker for the drafting endpoint; no external browser trackers. Model requests contain only a selected fictional task and closed-list drafting context, never candidate answers or names. Demo records remain on the device.

6. Template and human verification

Fictional task: reconcile a small shop's cash drawer. Opening cash MXN 500 + cash sales 1,890 - cash refund 120 - cash expenses 250 - bank deposit 1,000 = closing cash MXN 1,020. Changed condition: the refund was paid by card, not cash; corrected closing cash = MXN 1,140. Reference amounts support reviewer preparation; they never issue a candidate score.

Criteria: trace each cash movement; identify missing/contradictory evidence; explain the changed-condition adjustment; state assumptions and AI assistance. Reviewer records observed/support-needed/not-observed for each criterion, a factual note and an explanation of limits. No weighted total. The proof names reviewer, task version, observation method, revision, assistance and date.

Benchmark: Riipen's institution/employer project scoping, matching and feedback workflow is the closest operating reference; Forage shows bounded simulations, sample tasks and model answers. Umbral's localization bet is Spanish-first, informal-experience recognition, short named-human review and candidate-controlled sharing. This does not establish employer preference over an existing work-sample test. Official references checked for the implementation: <https://www.riipen.com/educators> and <https://www.theforage.com/>.

7. Shadow clause and acceptance checks

No unpaid production work. The fictional assessment cannot be repurposed as a client deliverable. Any later real-work branch requires agreed pay, scope, supervision, acceptance rules and applicable labor protections. No permanent failure label or public unsuccessful attempt. Candidate sees results first and controls every recipient. Accessibility accommodations must preserve the capability observed rather than reward polished presentation.

Acceptance: cannot invite without reviewer/deadlines/scope confirmation; cannot submit without assistance disclosure; cannot finish review without each criterion and observed revision/explanation; cannot approve unresolved correction; cannot export proof without candidate approval and named-recipient authorization; revocation disables future recipient export; substantive review edits invalidate prior sharing; deletion removes the local record; recover gracefully if browser storage or AI is unavailable; keyboard/mobile usable.

8. Mechanical and persona checks

Test the high-consequence state/consent guards and the known arithmetic. Exercise the actual browser flow at desktop/mobile, including one correction and one revocation. Record any observed defect, correction and subsequent deployment honestly; do not manufacture a defect or deployment history. Use a clearly labeled synthetic persona for comprehension and privacy checks, with captures from the actual app, and separate those results from real user validation.

Proposed pilot: 3 businesses, 12 candidates, not achieved. Stop/redesign if fewer than 2/3 employers would pay to repeat, median review exceeds 10 minutes, or comparison against final-artifact-only evidence reveals no useful additional hiring information. Demo fixtures do not count toward those targets. Money owns real willingness-to-pay research; this operator slice only makes delivery observable.

9. Three-year vision and next evidence

If a bounded sample changes actual employer decisions and reviewer costs stay manageable, Umbral could become a small institutional evidence service with reusable task families and interoperable candidate-owned proof. It would need authenticated privacy, reviewer calibration, retention controls, accessible observation, payment/repeat-purchase evidence and local partner capacity. Do not build a marketplace before that evidence.

Next evidence: compare proof versus final artifact with actual employers; observe informal-experience candidates completing the task; measure reviewer minutes; test willingness to fund repeat assessments; review candidate comprehension of recipient-specific consent. Team confirmation, videos, native build-chat export, public GitHub and a verified live LLM call remain separate delivery items.

10. Implementation update - 7 October 2026

The bounded operator workflow is implemented and its first private publication succeeded. Existing regression checks passed 14/14 after a real malformed-restored-review defect was corrected. An actual fictional browser walkthrough exercised candidate-first review, separate permission, revocation and correction followed by new private confirmation. Manual demo timer inputs are not measured pilot results.


The first export path could not confirm a downloaded file in the integrated browser. It now exposes visible JSON with text selection and manual saving instructions, plus a printable proof in the current app. Native browser download and OS Save as PDF remain unverified. Actual captures and the synthetic-persona report document the limits.

The Gemini drafting endpoint is wired, but the live provider secret remains missing; the UI explicitly says AI is unavailable. Production authentication, public access/GitHub, actual videos, original build-chat export and real pilot validation are not claimed complete. Published URL:

<https://umbral-week9-rodrigo.j6x567qt8g.chatgpt.site>.

11. Actual interface and workflow

Capture from the real local build. The 446-pixel viewport is the integrated browser default. Fictional states shown here are produced by the walkthrough; they are not pilot results.



A usable first rung needs an operating process

Named reviewer and deadlines precede the sample. The three fixture cases are labeled fictional. The ten-minute card is a proposed pilot threshold.

Privacy gates

Invitation > explanation > human review > private candidate result > accuracy confirmation > recipient permission.

Correction or substantive new review invalidates earlier approval and permission. An incomplete observation stays private and may be retried.

Scope of proof

One fictional task with version, reviewer, observed revision, assistance disclosure and a named recipient. No universal score or promise of employment.

12. Implementation specification

This is the implementation specification recorded for the current Codex build. It is not a raw export of the user conversation.

Build Umbral, the OPERATOR contribution to the supplied Team 2 First-Proof Blueprint. Preserve the Blueprint's draft status and Rodrigo's final individual brief. Store the Packet before product code. Use one bounded fictional cash-reconciliation task, not a CV builder or job board.

Implement operator intake with assigned reviewer and deadlines; task/rubric disclosure; assistance disclosure; human observation of a changed condition; evidence per criterion; candidate-first private result; correction; separate accuracy confirmation and named-recipient consent; revocation; private retry; portable proof export. Never rank, automatically reject, score universal employability or turn unsuccessful attempts into public failure labels. Do not create unpaid client deliverables.

Use semantic HTML/CSS, browser ES modules, versioned synthetic JSON and deterministic transitions. Keep device-local demo records plainly labeled. The role switch is exploratory navigation. Never imply production authentication.

Use a Worker endpoint for optional Gemini scope wording. Send only a closed fictional task ID and accommodation choice; reject candidate data and unexpected keys. Keep the provider secret on the server. Explicitly show unavailable AI when no secret is configured. Never present template text as a live LLM response.

Save actual commits, failure/correction observations, source provenance and deployments. Generate Packet and clearly synthetic Persona reports with captures from the actual application. Preserve the real build conversation separately when native export is available; do not invent a raw transcript. Prepare recording scripts without claiming that videos or interviews happened.

13. Delivery provenance

The Packet was first committed as c3a13ed before product-source code. Version 1 was published successfully from 91adef9722b03b9f083eeb818482e2a2ecc692b3. The managed source preserves meaningful commits. Deployment status and subsequent refinements are recorded in docs/DEPLOYMENT_LOG.md. The supplied Blueprint is preserved outside the app repository.

Sources: Team 2 Blueprint supplied by Rodrigo; Rodrigo's final individual Brain brief; actual app and Git history. Operating references: Riipen Educators (<https://www.riipen.com/educators>), Forage (<https://www.theforage.com/>) and official Gemini text-generation documentation (<https://ai.google.dev/gemini-api/docs/text-generation>). These references do not validate the proposed pilot.