

UMBRAL

Synthetic Persona Report / Candidate control

Generated roleplay and actual fictional browser observations. No participant interview, real human assessment, employer decision or pilot result.

Rodrigo Peña León - OPERATOR - Week 9 - 7 October 2026

1. Method and evidence boundary

This is a Codex-generated synthetic persona exercise using the actual Umbral interface and fictional records. No candidate or employer was interviewed. Codex operated the candidate and reviewer views; the reviewer notes are simulated, not a real human assessment. The exercise checks workflow behavior and interpretation risks. It cannot validate demand, fairness, accessibility or hiring outcomes.

The prompt is preserved in docs/PERSONA_PROMPT.md. Captures come from the actual local app at <http://127.0.0.1:4179>, not invented mockups. The browser viewport was 446 pixels wide; a requested desktop override did not change the measured DOM width, so a desktop-size pass is not claimed.

2. Invented persona and failed moment

Ruth, 24, has helped a family shop reconcile cash but lacks a formal credential. She can explain her procedure in writing and uses a spreadsheet. She wants to demonstrate one skill without losing control of mistaken feedback or exposing a private incomplete attempt. This profile is an invented stress case, not a claim about Mexican candidates.

Operator counterpart: a career-office coordinator with little reviewer capacity. The coordinator needs a named reviewer and date, usable notes, an observable revision and a measured review cost. A complete-looking proof alone does not show an employer will fund it.

3. Synthetic dialogue and interpretation

The following exchange is generated roleplay, not a participant transcript.

Ruth: "Who will read this, and can I explain it in writing?"

Facilitator interpretation: the invitation displays the assigned fictional reviewer, delivery/review dates, rubric and a written explanation option. This supports orientation; a real user comprehension check is still needed.

Ruth: "If I press Confirmar y mantener privado, have I sent it to a shop?"

Facilitator interpretation: no. Accuracy confirmation retains a private state. Recipient authorization requires a separate form with name, purpose, expiry and consent. The demo does not transmit files to the recipient.

Ruth: "The note should say this was a written explanation using fictional movements."

Observed action: Codex requested that correction. The UI displayed Corrección solicitada, removed the approval/share path and kept the old recipient revoked. After a simulated reviewer amended the explanation, the UI returned Resultado privado and required a new candidate confirmation.

Ruth: "Can I take the copy back after someone downloads it?"

Observed action: revocation removed the JSON/print controls for that recipient. The interface explicitly states downloaded copies cannot be recalled. This is a local future-export restriction, not remote deletion or cryptographic access control.

Ruth: "Does the record say I am employable?"

Observed UI: the export says the case is fictional and does not establish identity, employability or hiring. The rubric describes evidence in this task and produces no overall score. A real reviewer could still misuse a document outside the app; this exercise cannot establish that external behavior.

4. Actual browser observations

A. Review and candidate-first result. Codex filled all four evidence notes, recorded written observation and confirmed the changed condition. Saving returned to the candidate view with a private result. A manually entered seven-minute value tested the timer field; it is not a measured human-review median.

B. Accuracy and permission. Codex confirmed accuracy, then authorized the fictional Tienda Sol recipient for a stated purpose and 30-day period. The view showed the recipient, purpose, expiry and separate export/revoke controls.

C. Export. The generated JSON was visible and contained the task version, reviewer, observed revision, assistance disclosure and exact recipient authorization. Its displayed contents were saved as evidence/FICTIONAL_PROOF_UM-001.json. The embedded browser's download event and file-link download timed out; a successful native browser download is not claimed. The visible text selection fallback was checked. The printable proof preview was observed; an OS Save as PDF file was not verified.

D. Revocation and correction. Revocation removed export controls. The correction request kept the record private. A subsequent simulated review removed pending correction and required new confirmation, with the old permission still revoked.

E. Missing AI secret. The real scope-assistant button showed that AI was not configured and the reviewed template remained available. No live provider call or generated scope is claimed.

F. Narrow layout. The actual 446-pixel view stacked the sessions, reviewer facts, task and permission rules without horizontal overflow. This is a narrow-browser observation; it is not a full device/accessibility study.

5. Product changes prompted by the exercise

The initial export path automatically clicked a temporary link and showed a downloaded toast. The integrated browser did not produce a confirmed download. The revised path displays the exact JSON, exposes a download link and offers text selection with manual saving instructions. Toasts now say the copy is ready, not downloaded. Print preview stays in the current app rather than depending on a popup.

Separately, the mechanical checks found a malformed restored review that could crash the result screen. A strengthened restore validator now rejects that structure and shows a recovery message with new fictional cases. The passing checks demonstrate those tested rules, not production security.

6. What changed the operator's view

Generated interpretation, not Rodrigo's personal testimony: a polished proof is insufficient unless review capacity and consent operations are viable. Corrections create more work and invalidate previous permission. The software can make that work visible; it cannot establish demand or create a paid first rung by itself. This preserves the individual brief's concern about income and paid Micro-experience.

The proposed three-business/twelve-candidate pilot remains unperformed. Stop/redesign if fewer than two of three would pay to repeat, median review exceeds ten minutes, or observed explanations add no useful hiring information over the artifact alone.

7. Next real validation

Observe candidates explaining the actual task and consent choices; compare artifact-only and observed-proof decisions with employers; measure true reviewer time including corrections; check willingness to fund repetition; test written alternatives and assistive technology; verify export in standard browsers. These are next evidence tasks, not results in this report.

8. Interface evidence

Actual app captures. The correction request kept the case private; the AI helper explicitly reported missing configuration. Full captures and the visible fictional JSON are retained locally with the source.

Observed UI: Corrección solicitada. El registro sigue privado hasta resolverla.

Action performed with fictional data	Visible result
Request a correction to observation wording.	Approval/share path removed; correction pending; old recipient revoked.
Save the amended simulated review.	Resultado privado; unchecked accuracy confirmation required again.
Revoke a recipient permission.	Revocado; recipient export controls removed. Previously downloaded copies cannot be recalled.
Open scope wording assistance without provider secret.	La asistencia con IA todavía no está configurada. Puedes usar la plantilla revisada.
Prepare authorized JSON.	Visible task, reviewer, observed revision and exact recipient. Selection fallback worked; download completion unverified.

Capture provenance: evidence/correction-private.jpg, evidence/candidate-private.jpg, evidence/ai-unavailable.jpg and evidence/proof-preview.jpg. Full original images are kept with the local build; these text excerpts report observed labels without reducing long phone captures to unreadable figures.

Reproducibility and limits

The persona prompt, interpretation, browser observations and changes are preserved in docs/PERSONA_PROMPT.md, docs/PERSONA_REPORT.md and docs/TEST_LOG.md. Download completion and OS PDF saving remain unverified. The synthetic scenario supplies a test case, not evidence of user prevalence, institutional adoption or ability.